PORTFOLIO CASE STUDY

Agentic AI Staff

A 24-Bot Grok Bot Team Built to Scale a Book Series and a Business

Oct 10 – Nov 2, 2026

24	50	13 of 25	12
Bots on the Team	Work Tasks Closed on One Shared List	Book 2 Chapters Line-Edited in 13 Days	Site Pull Requests Merged in 4 Days

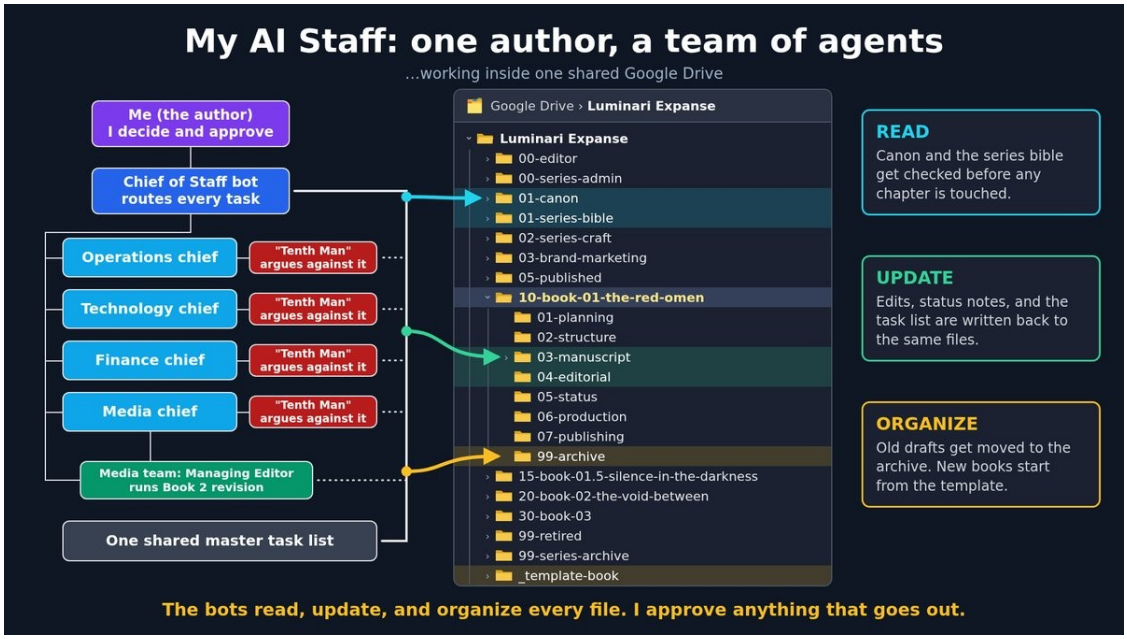
Section 1: What Was Built

Agentic AI Staff is a team of 24 Grok Bot bots that runs Daniel Stroup's author business and his AI LABORATORY™ work the way a staff runs an organization. It supports *The Luminari Expanse* science fiction series, the luminariexpanse.com author site, and the AI LABORATORY™ consulting site.

Every task goes to one bot, the Chief of Staff, which routes it to the right section. Under it are four section chiefs: Operations, Technology, Finance, and Media. Operations has three sub-leads: Publishing, Lab, and Business. Media's team is the Managing Editor, Marketing, the Publicist, and Audience. Technology works alongside the Grok Bot Platform Expert and the GitHub Engineer, and Finance draws on the existing Finance bot.

Next to every chief sits a Tenth Man: a designated dissenter that argues against its chief's recommendation and reports to that chief, never directly to Daniel. Every bot reports into one shared Master Task List, kept by the Chief of Staff, and every bot works from one shared Google Drive. The files are the team's memory.

Core Capability
A standing team of AI agents with one front door, one shared task list, shared files as memory, and four layers of governance. The bots draft and Daniel approves. The book-writing system and the website behind it were rebuilt so the team can keep growing.



The team: one Chief of Staff, four section chiefs each paired with a Tenth Man, one task list, and the shared Luminari Expanse Drive folder tree.

Section 2: Why It Was Built

I started using Grok Bot to upgrade how I wrote my books. I was managing 12+ separate chats, each scoped to one task or project, and that had become unwieldy. I also wanted to keep expanding my AI learning.

After some setup, Grok Bot proved more effective overall than the Claude setup I had been using, mainly because of connectors and task delegation. It also forced a redesign of how I run projects. Book writing and website updates had been slow and cumbersome as separate efforts.

The breadth of connectors lets the bots update many platforms and pull data from many systems, which makes the Chief of Staff better informed and more responsive. The built-in connector permissions and restrictions made the results easier to trust, and the bots ask before important tasks.

> **Core Question**
> Can one author replace a dozen separate, task-scoped AI chats with a coordinated team that runs several workstreams at once, with limited supervision, and still trust the results?

Section 3: Starting Competence

Grok Bot Platform	New. One of the first bots created was a Platform Expert to advise on the platform.
Prior AI Setup	Claude, used across 12+ separate chats, each scoped to one task or project
Prior Multi-Agent Work	A CrewAI multi-agent pipeline (an earlier AI LABORATORY™ case study), but never a standing team of bots coordinating with each other
Work Carried In	*The Luminari Expanse* series, including Book 2, *The Void Between*, plus luminariexpanse.com and the AI LABORATORY™ site

The team was designed and run by someone new to the platform.

Section 4: How AI Was Used

One Front Door

The Chief of Staff is the single point of entry. Daniel brings it a task, it routes the task to the right section, and it reports back with one voice. Other bots send their to-dos and status changes to the shared Master Task List instead of keeping their own lists. Before Oct 24, tasks were tracked in chat. On Oct 24 the list moved into one shared Google Sheet.

Each task on the list has an owner, a status, and a due date. A task without a date either gets one or gets set aside.

Section Chiefs and the Tenth Man

Daniel approved the four section chiefs' standing prompts on Oct 24. Each chief files one weekly report of about 400 words under four headings: Moving, Stalled, Needs your decision, and Opportunities. The reports are due Saturday at 9:30 AM MT for Daniel's 10 AM review with the Chief of Staff.

The four Tenth Man bots were created the same day at Daniel's request, one for each chief. Each one's job is to argue against its chief's recommendation. It reports to the chief, not to Daniel, so a recommendation has already been challenged before it reaches him.

> **Key Workflow Insight**
> The dissent happens inside the team. By the time a recommendation reaches Daniel, it has already been argued against, so he decides instead of debugging.

Rebuild 1: The Book-Writing System

On Oct 11, Book 2 was split into Books 2 and 3, with one folder per book. On Oct 25, the Luminari Expanse Drive went through a full file-structure revision. Book 2 and Book 3 folders were unified in

62 of 62 planned moves with 0 failures: 60 files moved, including 25 original drafts and 3 older drafts archived and 19 status files relocated, and 2 old folders retired. The same evening, 3 series-wide editor files were moved back to the editor folder, for 65 moves in total. Nothing was deleted. Bot file access is limited by design, so with Daniel's approval the moves ran through the desktop Drive folder.

The resulting layout has series folders, numbered book folders, a retired folder, and a template folder, so a new book starts from the template instead of from scratch. Bots read the canon and the series bible before touching a chapter, update files in place, and archive old drafts. The files are the shared memory.

The Managing Editor, created Oct 12, line-edits one chapter at a time. It posts numbered findings, Daniel rules on each item, and then it applies the approved items in one guarded write, verified byte for byte. Every ruling goes into a rulings log, which went from v2.22 (Oct 10–11) to v2.29 (Nov 1), 8 versions in the window.

Rebuild 2: The AI LABORATORY™ Site

The GitHub Engineer bot owns the AI LABORATORY™ site repository. On Oct 28 it rebuilt the site on Astro 7 with static output, one Markdown file per case study under a strict schema, and build and security checks, so a broken entry cannot reach the live site. The rebuild went live Nov 1, and deploys moved to GitHub Actions. A category-based portfolio followed the same day, with a library, category pages, and related projects.

Other merged work included consulting copy, trademark wording, copy edits, a new hero logo, and browse text. Nothing merges without Daniel's approval. On Nov 1 a Case Study Writer bot was added under the Chief of Technology, along with a case study kit: a structure guide, a starter template, and an inventory of every published study.

On luminariexpanse.com, on Oct 29, the team authenticated the email sender domain and turned the welcome email back on with an author-domain sender. It also published Privacy Policy and Terms of Service pages, fixed the /free/ cover image and share preview, removed placeholder contact details from the footer, and set 2 test pages to Draft.

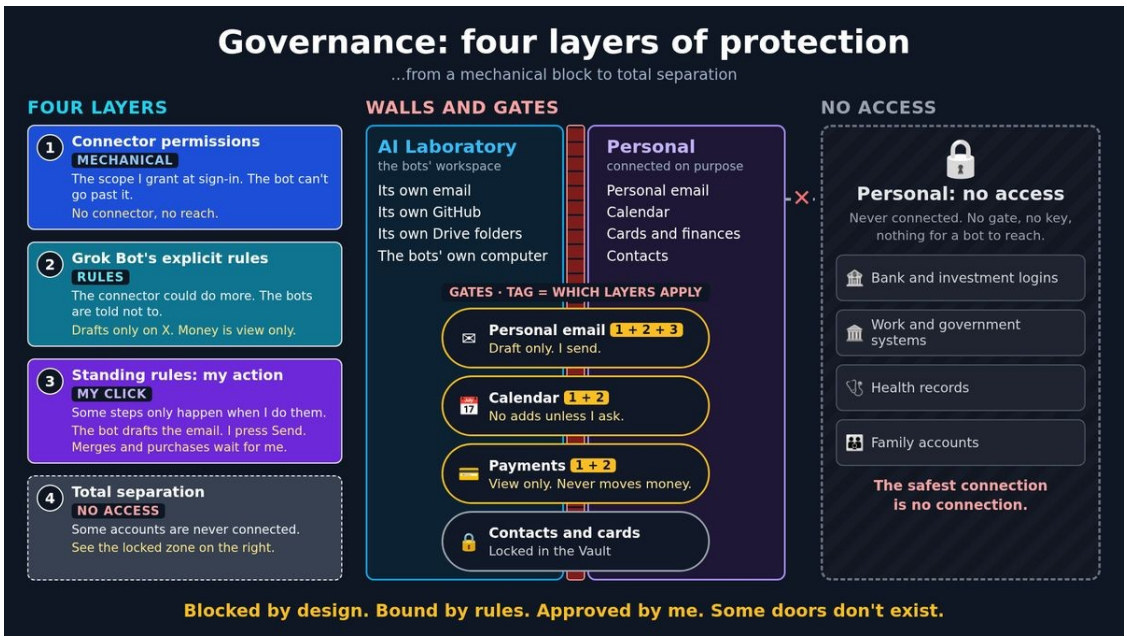
> **Key Workflow Insight**
> Both rebuilds were built to scale. A new book starts from the template folder, a new case study is one new file checked by the build, and new bots plug into the same front door and the same task list.

Governance: Four Layers

- **Connector permissions.** The scope granted at sign-in. A bot cannot go past it.
- **Explicit rules.** A connector could allow more, and the bots are told not to.
- **Steps that wait for Daniel's click.** The bot drafts the email and Daniel presses Send. Merges and purchases wait for him.
- **Total separation.** Some accounts are never connected.

The bots work in an AI LABORATORY™ workspace walled off from personal accounts, with gates opened only on purpose. Operations runs a data-security checklist on every Lab build, which

includes an AI data-flow statement and independent verification by a separate AI model. A weekly Platform Utilization Review, set up Oct 24 and refocused Oct 31, looks for platform features the team is not yet using.



Governance: connector permissions, explicit bot rules, steps that wait for Daniel's click, and total separation.

Models

Grok Bot ran every bot on the team; model versions not recorded. The earlier separate-chats setup used Claude; model versions not recorded.

Section 5: Timeline & Conditions

Oct 10, 2026	Chief of Staff's first daily brief
By Oct 11	First specialist bots created: Grok Bot Platform Expert, Finance, Research, Publicist, Marketing, Scheduler, Strategist
Oct 11	Book 2 split into Books 2 and 3, one folder per book
Oct 12	Managing Editor created; Book 2 first draft complete (Prologue, ch.1–23, Epilogue, about 112,700 words)
Oct 20	Book 2 line edit begins (Prologue and ch.1)
Oct 24	Section-chief prompts approved; four Tenth Man bots created; Master Task List becomes the single shared list; weekly Platform Utilization Review set up
Oct 25	Drive file-structure revision: 65 moves, 0 failed, nothing deleted
Oct 28	Astro rebuild of the AI LABORATORY™ site built
Oct 29	luminariexpanse.com fixes; first site pull requests merged
Nov 1	Astro rebuild live; deploys on GitHub Actions; category portfolio; Case Study Writer bot and case study kit; Book 2 line edit reaches ch.12
Total Duration	About three weeks, Oct 10 – Nov 2, 2026

Conditions	Built while the work ran: the Book 2 revision and the site work went on alongside the setup. The team works within a weekly usage limit.

Section 6: Outcome & Evidence

Since the Chief of Staff and the team were set up, Daniel's tally covers several areas. Chapters were edited and written, and the website got numerous updates, including a full architectural rebuild. The Luminari Expanse file structure was completely revised, and the way the website and its case studies are produced was redesigned. Each of these architectural overhauls was tied to the Grok Bot system's connectivity. Daniel can now run several workstreams at once and set large workloads in motion with limited supervision.

Counted Results

- **Master Task List** (read Nov 2, 2026): 119 work tasks, 50 done, across the four work sections: Media 16, Technology 15 plus 1 shared with the Chief of Staff, Finance 11, and Operations 7. In total, including family and personal rows, the list tracks 140 tasks, 53 done. The list grew from 91 rows with 18 done on Oct 24 to 130 rows with 31 done on Oct 28.
- **Book 2 line edit:** 13 of 25 chapters (Prologue through ch.12) line-edited between Oct 20 and Nov 1, including 4 chapters in the week to Nov 1.
- **Book 2 drafting:** the first draft was completed Oct 12, with ch.23 and the new Epilogue written inside the window.
- **AI LABORATORY™ site:** 12 pull requests merged Oct 29 – Nov 1, and 50 commits on the main branch in the window (34 not counting merges), including the full Astro rebuild.
- **Drive:** 65 file and folder moves, 0 failed, nothing deleted.
- **luminariexpanse.com:** sender domain authenticated, welcome email restored, legal pages published, free-book page fixed.

Section 7: Key Metrics

Bots on the Team	24
Structure	1 Chief of Staff · 4 section chiefs · 4 Tenth Man dissenters · specialists under each chief
Task System	1 shared Master Task List; 119 work tasks, 50 done, as of Nov 2, 2026 (140 in total including family and personal rows, 53 done)
Weekly Reporting	4 chief reports, about 400 words each, 4 headings, Saturday review
Book 2 Line Edit	13 of 25 chapters in 13 days (Oct 20 – Nov 1)
Rulings Log	v2.22 → v2.29, 8 versions
Drive Revision	65 moves, 0 failed, nothing deleted (Oct 25)
Site Pull Requests	12 merged in 4 days (Oct 29 – Nov 1)
Site Commits	50 on main (34 not counting merges)
Site Platform	Astro 7 static site, one file per case study, build and security checks, GitHub Actions deploys
Governance Layers	4: connector permissions, explicit rules, Daniel's click, total separation

Models	Grok Bot; earlier setup Claude. Model versions not recorded

Section 8: Lessons Learned

Lesson 1: You Don't Have to Design the System Yourself

Define the criteria, what success looks like, and the problems to avoid. Then ask the bot for recommendations, and make it ask you clarifying questions.

Lesson 2: Make It Read Everything Back Before It Builds

Before anything is implemented, have Grok Bot read everything back so you can confirm it understands the goal and has a verified structure for its approach.

Lesson 3: Ask for Pictures, Not Just Words

Have it produce visual representations so you can see the structure, not just read it.

Lesson 4: Trust Comes From the Guardrails

The built-in connector permissions and restrictions made the results easier to trust, and so did bots that ask before important tasks.

What This Project Proves

- **A staff, not a set of chats:** 24 bots coordinated through one front door and one task list, with 50 work tasks closed.
- **Built to scale:** the book-writing system and the AI LABORATORY™ site were both rebuilt so new books, case studies, and bots plug into the same structure.
- **Real throughput:** 13 of 25 Book 2 chapters line-edited in 13 days and 12 site pull requests merged in 4 days.
- **Governed by design:** four layers of protection; the bots draft and Daniel approves.

--- END OF CASE STUDY ---

ailaboratorydesk.com